SpotBot status report

Monday, October 26, 2026 · end of Phase 2, Phase 3 starts October 27 · prepared for the standup push

13
bugs fixed (9 from the audit, 4 found by the new tests)

1,167
offline tests passing, 0 failing (was 995)

10
local commits ready to push on Oct 27

580
Phase 3 points ready to seed on the sprint board

What changed today

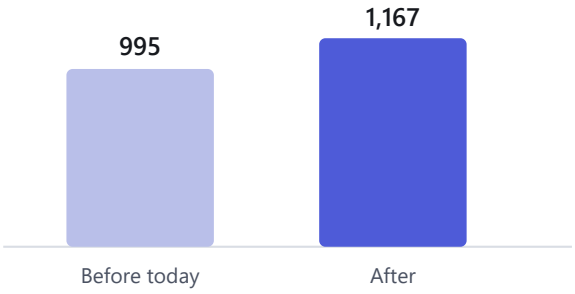
- **The "Not a place" list can no longer lose a spot.** A moved spot is the only copy left, so it is now saved under its own id, kept past the 30-day expiry and the 300-item cap, and only removed after "Move it back" succeeds.
- **Bot cards behave under load.** Buttons answer inside Discord's 3-second limit, wait past the model's 90-second timeout, and failed links next to not-a-place cards get Retry / Add manually.
- **The Instagram webhook never returns a 500.** Any delivery shape is skipped safely (Meta resends every non-200), and every sender id is hashed in the event log.
- **Sprint board gets an `archive` command.** It writes one index post that links every Phase 2 standup, digest and finished feature, then closes those posts. Nothing is deleted.

Bugs fixed

Area	Bug	Type
Not-a-place list	Two manual spots shared one key, so the second move overwrote the first	data loss
Not-a-place list	Moved spots expired after 30 days or fell off the 300 cap	data loss
Restore	Entry removed before the spot was re-added, so a failure lost both	data loss
Bot buttons	Move it back missed Discord's 3 s deadline; Not a place timed out before the model did	UX
Bot cards	A no-text reel was called "not a place"; missed links had no Retry button	UX
Post summaries	A reply cut off at the token limit showed raw JSON on the card	UX
Post summaries *(new)*	Non-text titles and 600-character "two-word" titles reached cards	UX
Webhook	Unexpected payload shapes raised a 500 and caused endless redelivery	robustness
Webhook *(new)*	Very deeply nested JSON crashed the parser (500 → now 400)	robustness
Event log *(new)*	A sender id sent as plain text was logged unhashed	privacy

13 fixes in total. Rows group related fixes. *(new)* means the bug was found by the new adversarial tests, not the audit.

Tests

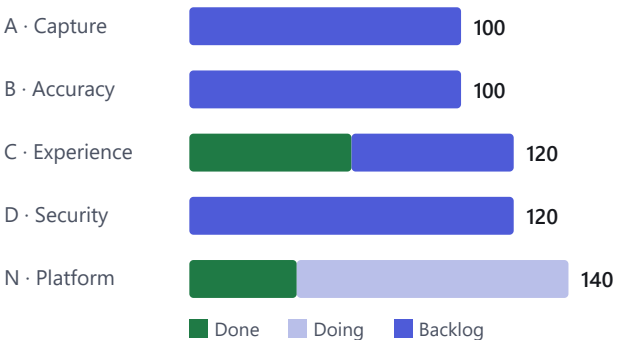


Offline suite (`pytest -k "not live"`): 0 failures, 8 expected failures (xfail).

New adversarial tests (`test_adversarial.py`, seeded so they repeat):

- 1,600 random and mutated webhook deliveries
- Random move / paste / expire / restore sequences, checking a moved spot is always recoverable
- 8 threads writing the list at once
- A model reply cut off at every possible character

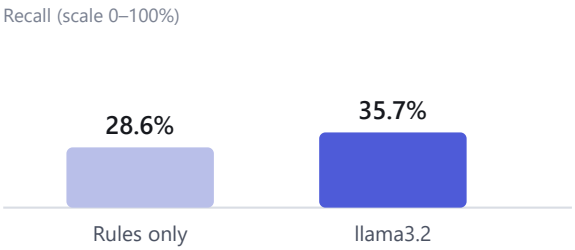
Phase 3 on the sprint board



8 features, 580 points. Starting status comes from each plan's checkboxes.

Feature	Pts	Start
#15 Gatekept-reel recovery ★	100	Backlog
#6 Place-or-not gate ★	100	Backlog
#21 Privacy & data-deletion UX ★	60	Done
#22 Usability fixes from the study	60	Backlog
#13 Threat model & security review	60	Backlog
#31 Dependency & supply-chain scanning	60	Backlog
#11 Always-on staging deployment	100	Doing
#26 Retry only what's worth retrying	40	Done

Accuracy: the not-a-place detector (dev split)



The LLM catches 5 of 14 not-a-place reels against 4 for the rules. Its precision is 71.4% (the rules got all 4 of their picks right). With only 14 examples the gap is small and uncertain. A larger model (llama3.1:8b) is being scored next. The test split has not been read.

Next

- **Oct 27:** push the 10 commits for standup. Then, from the session with the SprintBot token: `archive --phase 3 --apply`, then `seed --phase 3 --apply`.
- **Voting card:** the top pick as a poster/ticket hero with "+1" approval votes, a private "Show more" list, auto-close at quorum, and a random tie-breaker with a spin animation.
- **/prefs:** people enter their own preferences, which feed `/plan`. The bot does not listen to chat.
- **Housekeeping:** move root tests into `tests/`, sort the docs, rewrite the README (it still says 143 tests).